

Lab 8 Assignment

July 30, 2025

1 Lab 8: Define and Solve an ML Problem of Your Choosing

```
[1]: import pandas as pd
import numpy as np
import os
import matplotlib.pyplot as plt
import seaborn as sns
```

In this lab assignment, you will follow the machine learning life cycle and implement a model to solve a machine learning problem of your choosing. You will select a data set and choose a predictive problem that the data set supports. You will then inspect the data with your problem in mind and begin to formulate a project plan. You will then implement the machine learning project plan.

You will complete the following tasks:

1. Build Your DataFrame
2. Define Your ML Problem
3. Perform exploratory data analysis to understand your data.
4. Define Your Project Plan
5. Implement Your Project Plan:
 - Prepare your data for your model.
 - Fit your model to the training data and evaluate your model.
 - Improve your model's performance.

1.1 Part 1: Build Your DataFrame

You will have the option to choose one of four data sets that you have worked with in this program:

- The "census" data set that contains Census information from 1994: `censusData.csv`
- Airbnb NYC "listings" data set: `airbnbListingsData.csv`
- World Happiness Report (WHR) data set: `WHR2018Chapter20onlineData.csv`
- Book Review data set: `bookReviewsData.csv`

Note that these are variations of the data sets that you have worked with in this program. For example, some do not include some of the preprocessing necessary for specific models.

Load a Data Set and Save it as a Pandas DataFrame The code cell below contains filenames (path + filename) for each of the four data sets available to you.

Task: In the code cell below, use the same method you have been using to load the data using `pd.read_csv()` and save it to DataFrame `df`.

You can load each file as a new DataFrame to inspect the data before choosing your data set.

```
[2]: import os
import pandas as pd
# File names of the four data sets
adultDataSet_filename = os.path.join(os.getcwd(), "data", "censusData.csv")
airbnbDataSet_filename = os.path.join(os.getcwd(), "data", "airbnbListingsData.
↪csv")
WHRDataSet_filename = os.path.join(os.getcwd(), "data", "WHR2018Chapter2OnlineData.csv")
bookReviewDataSet_filename = os.path.join(os.getcwd(), "data", "bookReviewsData.
↪csv")

df = pd.read_csv(airbnbDataSet_filename)

df.head()
```

```
[2]:
```

	name \		description \	neighborhood_overview	host_name \	host_location \
0	Skylit Midtown Castle		Beautiful, spacious skylit studio in the heart...	Centrally located in the heart of Manhattan ju...	Jennifer	New York, New York, United States
1	Whole flr w/private bdrm, bath & kitchen(pls r...		Enjoy 500 s.f. top floor in 1899 brownstone, w...	Just the right mix of urban center and local n...	LisaRoxanne	New York, New York, United States
2	Spacious Brooklyn Duplex, Patio + Garden		We welcome you to stay in our lovely 2 br dupl...		NaN	Brooklyn, New York, United States
3	Large Furnished Room Near B'way		Please don't expect the luxury here just a bas...	Theater district, many restaurants around here.	Shunichi	New York, New York, United States
4	Cozy Clean Guest Room - Family Apt		Our best guests are seeking a safe, clean, spa...	Our neighborhood is full of restaurants and ca...	MaryEllen	New York, New York, United States

	host_about	host_response_rate	\
0	A New Yorker since 2000! My passion is creatin...	0.80	
1	Laid-back Native New Yorker (formerly bi-coast...	0.09	
2	Rebecca is an artist/designer, and Henoch is i...	1.00	
3	I used to work for a financial industry but no...	1.00	
4	Welcome to family life with my oldest two away...	NaN	

	host_acceptance_rate	host_is_superhost	host_listings_count	...	\
0	0.17	True	8.0	...	
1	0.69	True	1.0	...	
2	0.25	True	1.0	...	
3	1.00	True	1.0	...	
4	NaN	True	1.0	...	

	review_scores_communication	review_scores_location	review_scores_value	\
0	4.79	4.86	4.41	
1	4.80	4.71	4.64	
2	5.00	4.50	5.00	
3	4.42	4.87	4.36	
4	4.95	4.94	4.92	

	instant_bookable	calculated_host_listings_count	\
0	False	3	
1	False	1	
2	False	1	
3	False	1	
4	False	1	

	calculated_host_listings_count_entire_homes	\
0	3	
1	1	
2	1	
3	0	
4	0	

	calculated_host_listings_count_private_rooms	\
0	0	
1	0	
2	0	
3	1	
4	1	

	calculated_host_listings_count_shared_rooms	reviews_per_month	\
0	0	0.33	
1	0	4.86	
2	0	0.02	

3	0	3.68
4	0	0.87

	n_host_verifications
0	9
1	6
2	3
3	4
4	7

[5 rows x 50 columns]

1.2 Part 2: Define Your ML Problem

Next you will formulate your ML Problem. In the markdown cell below, answer the following questions:

1. List the data set you have chosen.
2. What will you be predicting? What is the label?
3. Is this a supervised or unsupervised learning problem? Is this a clustering, classification or regression problem? Is it a binary classification or multi-class classification problem?
4. What are your features? (note: this list may change after you explore your data)
5. Explain why this is an important problem. In other words, how would a company create value with a model that predicts this label?

Dataset Chosen: Airbnb NYC Listings (airbnbListingsData.csv)

Prediction Goal / Label: I will be predicting the price of a listing per night.

Type of ML Problem:

This is a supervised learning problem.

It is a regression problem because the label (price) is a continuous numeric value.

Initial Feature Set (subject to change after EDA):

neighbourhood_group (e.g., Manhattan, Brooklyn)

neighbourhood

room_type

minimum_nights

number_of_reviews

reviews_per_month

availability_365

latitude, longitude

1.3 Part 3: Understand Your Data

The next step is to perform exploratory data analysis. Inspect and analyze your data set with your machine learning problem in mind. Consider the following as you inspect your data:

1. What data preparation techniques would you like to use? These data preparation techniques may include:
 - addressing missingness, such as replacing missing values with means
 - finding and replacing outliers
 - renaming features and labels
 - finding and replacing outliers
 - performing feature engineering techniques such as one-hot encoding on categorical features
 - selecting appropriate features and removing irrelevant features
 - performing specific data cleaning and preprocessing techniques for an NLP problem
 - addressing class imbalance in your data sample to promote fair AI
2. What machine learning model (or models) you would like to use that is suitable for your predictive problem and data?
 - Are there other data preparation techniques that you will need to apply to build a balanced modeling data set for your problem and model? For example, will you need to scale your data?
3. How will you evaluate and improve the model's performance?
 - Are there specific evaluation metrics and methods that are appropriate for your model?

Think of the different techniques you have used to inspect and analyze your data in this course. These include using Pandas to apply data filters, using the Pandas `describe()` method to get insight into key statistics for each column, using the Pandas `dtypes` property to inspect the data type of each column, and using Matplotlib and Seaborn to detect outliers and visualize relationships between features and labels. If you are working on a classification problem, use techniques you have learned to determine if there is class imbalance.

Task: Use the techniques you have learned in this course to inspect and analyze your data. You can import additional packages that you have used in this course that you will need to perform this task.

Note: You can add code cells if needed by going to the Insert menu and clicking on Insert Cell Below in the drop-down menu.

```
[3]: # Import required libraries
import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Load the Airbnb dataset
```

```

airbnbDataSet_filename = os.path.join(os.getcwd(), "data", "airbnbListingsData.
↪csv")
df = pd.read_csv(airbnbDataSet_filename)

# Display basic info
print("Shape of dataset:", df.shape)
display(df.head())

# Check for missing values
print("\nMissing Values:\n", df.isnull().sum())

# Check data types
print("\nData Types:\n", df.dtypes)

# Summary statistics
print("\nDescriptive Stats:\n")
display(df.describe())

# Check number of unique values in each column
print("\nUnique Values:\n", df.nunique())

# Check price distribution
plt.figure(figsize=(8, 4))
sns.histplot(df['price'], bins=100, kde=True)
plt.xlim(0, 500) # Clip outliers for better visualization
plt.title("Price Distribution")
plt.show()

# Check for outliers in minimum_nights
plt.figure(figsize=(8, 2))
sns.boxplot(x=df['minimum_nights'])
plt.xlim(0, 100) # Clip for visualization
plt.title("Minimum Nights")
plt.show()

# Boxplot: Price by room_type
plt.figure(figsize=(6, 4))
sns.boxplot(x='room_type', y='price', data=df)
plt.ylim(0, 500)
plt.title("Price by Room Type")
plt.show()

# Correlation heatmap
plt.figure(figsize=(10, 6))
corr = df.corr(numeric_only=True)
sns.heatmap(corr, annot=True, cmap='coolwarm')
plt.title("Correlation Heatmap")

```

```

plt.show()

# Basic cleaning steps:

# Fill missing reviews_per_month with 0
df['reviews_per_month'] = df['reviews_per_month'].fillna(0)

# Drop unnecessary columns
df.drop(columns=['name', 'host_name', 'last_review', 'id'], inplace=True)

# Remove listings with extreme prices and minimum_nights
df = df[(df['price'] <= 500) & (df['minimum_nights'] <= 365)]

# One-hot encode categorical features
df = pd.get_dummies(df, columns=['neighbourhood_group', 'room_type'],
→drop_first=True)

# Confirm final structure
print("\nCleaned DataFrame Shape:", df.shape)
display(df.head())

```

Shape of dataset: (28022, 50)

	name \
0	Skylit Midtown Castle
1	Whole flr w/private bdrm, bath & kitchen(pls r...
2	Spacious Brooklyn Duplex, Patio + Garden
3	Large Furnished Room Near B'way
4	Cozy Clean Guest Room - Family Apt

	description \
0	Beautiful, spacious skylit studio in the heart...
1	Enjoy 500 s.f. top floor in 1899 brownstone, w...
2	We welcome you to stay in our lovely 2 br dupl...
3	Please don't expect the luxury here just a bas...
4	Our best guests are seeking a safe, clean, spa...

	neighborhood_overview	host_name \
0	Centrally located in the heart of Manhattan ju...	Jennifer
1	Just the right mix of urban center and local n...	LisaRoxanne
2	NaN	Rebecca
3	Theater district, many restaurants around here.	Shunichi
4	Our neighborhood is full of restaurants and ca...	MaryEllen

	host_location \
0	New York, New York, United States
1	New York, New York, United States
2	Brooklyn, New York, United States

3 New York, New York, United States
 4 New York, New York, United States

	host_about	host_response_rate	\
0	A New Yorker since 2000! My passion is creatin...	0.80	
1	Laid-back Native New Yorker (formerly bi-coast...	0.09	
2	Rebecca is an artist/designer, and Henoch is i...	1.00	
3	I used to work for a financial industry but no...	1.00	
4	Welcome to family life with my oldest two away...	NaN	

	host_acceptance_rate	host_is_superhost	host_listings_count	...	\
0	0.17	True	8.0	...	
1	0.69	True	1.0	...	
2	0.25	True	1.0	...	
3	1.00	True	1.0	...	
4	NaN	True	1.0	...	

	review_scores_communication	review_scores_location	review_scores_value	\
0	4.79	4.86	4.41	
1	4.80	4.71	4.64	
2	5.00	4.50	5.00	
3	4.42	4.87	4.36	
4	4.95	4.94	4.92	

	instant_bookable	calculated_host_listings_count	\
0	False	3	
1	False	1	
2	False	1	
3	False	1	
4	False	1	

	calculated_host_listings_count_entire_homes	\
0	3	
1	1	
2	1	
3	0	
4	0	

	calculated_host_listings_count_private_rooms	\
0	0	
1	0	
2	0	
3	1	
4	1	

	calculated_host_listings_count_shared_rooms	reviews_per_month	\
0	0	0.33	
1	0	4.86	

2	0	0.02
3	0	3.68
4	0	0.87

n_host_verifications	
0	9
1	6
2	3
3	4
4	7

[5 rows x 50 columns]

Missing Values:

name	5
description	570
neighborhood_overview	9816
host_name	0
host_location	60
host_about	10945
host_response_rate	11843
host_acceptance_rate	11113
host_is_superhost	0
host_listings_count	0
host_total_listings_count	0
host_has_profile_pic	0
host_identity_verified	0
neighbourhood_group_cleansed	0
room_type	0
accommodates	0
bathrooms	0
bedrooms	2918
beds	1354
amenities	0
price	0
minimum_nights	0
maximum_nights	0
minimum_minimum_nights	0
maximum_minimum_nights	0
minimum_maximum_nights	0
maximum_maximum_nights	0
minimum_nights_avg_ntm	0
maximum_nights_avg_ntm	0
has_availability	0
availability_30	0
availability_60	0

availability_90	0
availability_365	0
number_of_reviews	0
number_of_reviews_ltm	0
number_of_reviews_l30d	0
review_scores_rating	0
review_scores_cleanliness	0
review_scores_checkin	0
review_scores_communication	0
review_scores_location	0
review_scores_value	0
instant_bookable	0
calculated_host_listings_count	0
calculated_host_listings_count_entire_homes	0
calculated_host_listings_count_private_rooms	0
calculated_host_listings_count_shared_rooms	0
reviews_per_month	0
n_host_verifications	0
dtype: int64	

Data Types:

name	object
description	object
neighborhood_overview	object
host_name	object
host_location	object
host_about	object
host_response_rate	float64
host_acceptance_rate	float64
host_is_superhost	bool
host_listings_count	float64
host_total_listings_count	float64
host_has_profile_pic	bool
host_identity_verified	bool
neighbourhood_group_cleansed	object
room_type	object
accommodates	int64
bathrooms	float64
bedrooms	float64
beds	float64
amenities	object
price	float64
minimum_nights	int64
maximum_nights	int64
minimum_minimum_nights	float64
maximum_minimum_nights	float64
minimum_maximum_nights	float64
maximum_maximum_nights	float64

minimum_nights_avg_ntm	float64
maximum_nights_avg_ntm	float64
has_availability	bool
availability_30	int64
availability_60	int64
availability_90	int64
availability_365	int64
number_of_reviews	int64
number_of_reviews_ltm	int64
number_of_reviews_l30d	int64
review_scores_rating	float64
review_scores_cleanliness	float64
review_scores_checkin	float64
review_scores_communication	float64
review_scores_location	float64
review_scores_value	float64
instant_bookable	bool
calculated_host_listings_count	int64
calculated_host_listings_count_entire_homes	int64
calculated_host_listings_count_private_rooms	int64
calculated_host_listings_count_shared_rooms	int64
reviews_per_month	float64
n_host_verifications	int64
dtype:	object

Descriptive Stats:

	host_response_rate	host_acceptance_rate	host_listings_count	\
count	16179.000000	16909.000000	28022.000000	
mean	0.906901	0.791953	14.554778	
std	0.227282	0.276732	120.721287	
min	0.000000	0.000000	0.000000	
25%	0.940000	0.680000	1.000000	
50%	1.000000	0.910000	1.000000	
75%	1.000000	1.000000	3.000000	
max	1.000000	1.000000	3387.000000	

	host_total_listings_count	accommodates	bathrooms	bedrooms	\
count	28022.000000	28022.000000	28022.000000	25104.000000	
mean	14.554778	2.874491	1.142174	1.329708	
std	120.721287	1.860251	0.421132	0.700726	
min	0.000000	1.000000	0.000000	1.000000	
25%	1.000000	2.000000	1.000000	1.000000	
50%	1.000000	2.000000	1.000000	1.000000	
75%	3.000000	4.000000	1.000000	1.000000	
max	3387.000000	16.000000	8.000000	12.000000	

	beds	price	minimum_nights	...	review_scores_checkin	\
count	26668.000000	28022.000000	28022.000000	...	28022.000000	
mean	1.629556	154.228749	18.689387	...	4.814300	
std	1.097104	140.816605	25.569151	...	0.438603	
min	1.000000	29.000000	1.000000	...	0.000000	
25%	1.000000	70.000000	2.000000	...	4.810000	
50%	1.000000	115.000000	30.000000	...	4.960000	
75%	2.000000	180.000000	30.000000	...	5.000000	
max	21.000000	1000.000000	1250.000000	...	5.000000	

	review_scores_communication	review_scores_location	\
count	28022.000000	28022.000000	
mean	4.808041	4.750393	
std	0.464585	0.415717	
min	0.000000	0.000000	
25%	4.810000	4.670000	
50%	4.970000	4.880000	
75%	5.000000	5.000000	
max	5.000000	5.000000	

	review_scores_value	calculated_host_listings_count	\
count	28022.000000	28022.000000	
mean	4.647670	9.581900	
std	0.518023	32.227523	
min	0.000000	1.000000	
25%	4.550000	1.000000	
50%	4.780000	1.000000	
75%	5.000000	3.000000	
max	5.000000	421.000000	

	calculated_host_listings_count_entire_homes	\
count	28022.000000	
mean	5.562986	
std	26.121426	
min	0.000000	
25%	0.000000	
50%	1.000000	
75%	1.000000	
max	308.000000	

	calculated_host_listings_count_private_rooms	\
count	28022.000000	
mean	3.902077	
std	17.972386	
min	0.000000	
25%	0.000000	
50%	0.000000	
75%	1.000000	

max 359.000000

	calculated_host_listings_count_shared_rooms	reviews_per_month \
count	28022.000000	28022.000000
mean	0.048283	1.758325
std	0.442459	4.446143
min	0.000000	0.010000
25%	0.000000	0.130000
50%	0.000000	0.510000
75%	0.000000	1.830000
max	8.000000	141.000000


	n_host_verifications
count	28022.000000
mean	5.169510
std	2.028497
min	1.000000
25%	4.000000
50%	5.000000
75%	7.000000
max	13.000000

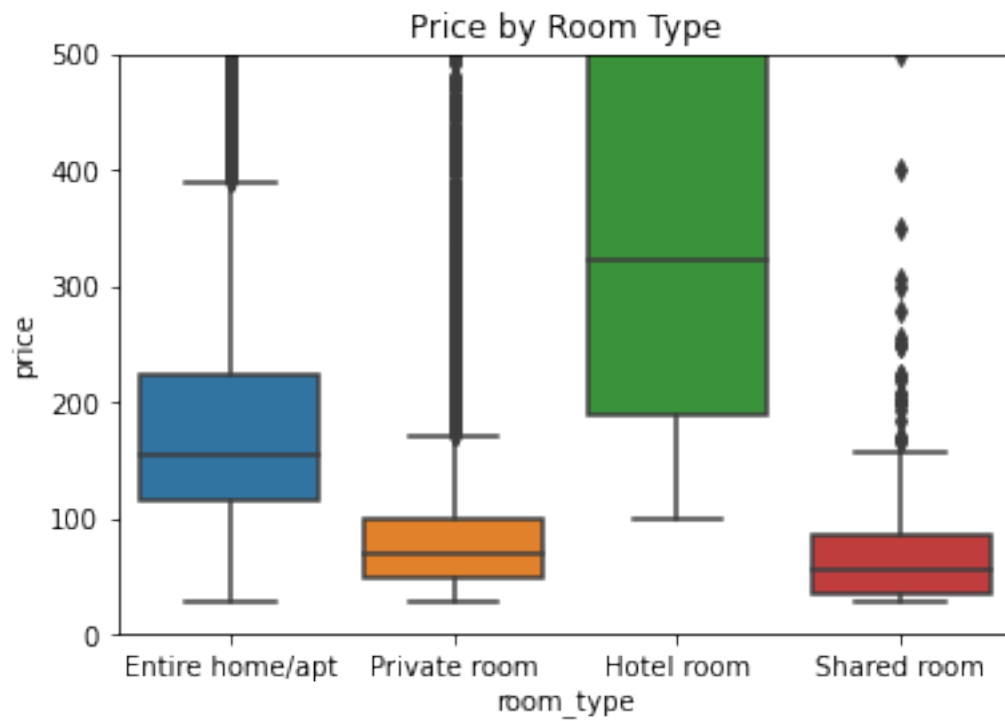

[8 rows x 36 columns]

Unique Values:

name	27386
description	25952
neighborhood_overview	15800
host_name	7566
host_location	1364
host_about	11962
host_response_rate	85
host_acceptance_rate	101
host_is_superhost	1
host_listings_count	73
host_total_listings_count	73
host_has_profile_pic	1
host_identity_verified	1
neighbourhood_group_cleansed	5
room_type	4
accommodates	16
bathrooms	16
bedrooms	11
beds	16
amenities	25020
price	684

minimum_nights	95
maximum_nights	229
minimum_minimum_nights	98
maximum_minimum_nights	102
minimum_maximum_nights	206
maximum_maximum_nights	206
minimum_nights_avg_ntm	329
maximum_nights_avg_ntm	452
has_availability	2
availability_30	31
availability_60	61
availability_90	91
availability_365	366
number_of_reviews	418
number_of_reviews_ltm	140
number_of_reviews_l30d	29
review_scores_rating	154
review_scores_cleanliness	196
review_scores_checkin	135
review_scores_communication	141
review_scores_location	153
review_scores_value	164
instant_bookable	2
calculated_host_listings_count	59
calculated_host_listings_count_entire_homes	44
calculated_host_listings_count_private_rooms	45
calculated_host_listings_count_shared_rooms	9
reviews_per_month	1357
n_host_verifications	13
dtype: int64	





↩ ⏏
↩ last) TypeError Traceback (most recent call ⏏

```

<ipython-input-3-0f846d227e34> in <module>()
    50 # Correlation heatmap
    51 plt.figure(figsize=(10, 6))
--> 52 corr = df.corr(numeric_only=True)
    53 sns.heatmap(corr, annot=True, cmap='coolwarm')
    54 plt.title("Correlation Heatmap")

```

TypeError: corr() got an unexpected keyword argument 'numeric_only'

<Figure size 720x432 with 0 Axes>

1.4 Part 4: Define Your Project Plan

Now that you understand your data, in the markdown cell below, define your plan to implement the remaining phases of the machine learning life cycle (data preparation, modeling, evaluation) to solve your ML problem. Answer the following questions:

- Do you have a new feature list? If so, what are the features that you chose to keep and remove after inspecting the data?
- Explain different data preparation techniques that you will use to prepare your data for modeling.
- What is your model (or models)?
- Describe your plan to train your model, analyze its performance and then improve the model. That is, describe your model building, validation and selection plan to produce a model that generalizes well to new data.

For this project, I chose to retain features that are most relevant to predicting Airbnb listing prices, including latitude, longitude, minimum_nights, number_of_reviews, reviews_per_month, and availability_365, along with one-hot encoded versions of neighbourhood_group and room_type. I removed non-predictive or high-cardinality columns such as id, name, host_name, last_review, and neighbourhood. To prepare the data for modeling, I handled missing values by filling reviews_per_month with 0, removed extreme outliers from price and minimum_nights, and applied one-hot encoding to categorical variables. My primary model will be a Random Forest Regressor due to its robustness with non-linear data and mixed feature types. I will also use a Linear Regression model as a baseline and may explore Gradient Boosting for further improvement. I plan to split the data into training and test sets, train and evaluate the models using RMSE, MAE, and R^2 metrics, and perform hyperparameter tuning through cross-validation to optimize model performance. My goal is to select a model that generalizes well to new data and provides actionable insights into what factors most impact Airbnb pricing.

1.5 Part 5: Implement Your Project Plan

Task: In the code cell below, import additional packages that you have used in this course that you will need to implement your project plan.


```
[4]: # Data manipulation and analysis
import pandas as pd
import numpy as np

# Data visualization
import matplotlib.pyplot as plt
import seaborn as sns

# Data preprocessing
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler # Optional for linear models
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Machine learning models
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor

# Model tuning (optional)
from sklearn.model_selection import GridSearchCV, cross_val_score

# Set plot style
sns.set(style='whitegrid')
```

Task: Use the rest of this notebook to carry out your project plan.

You will:

1. Prepare your data for your model.
2. Fit your model to the training data and evaluate your model.
3. Improve your model's performance by performing model selection and/or feature selection techniques to find best model for your problem.

Add code cells below and populate the notebook with commentary, code, analyses, results, and figures as you see fit.

```
[7]: # 1. Import all required packages
import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, GridSearchCV
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Set plotting style
```

```

sns.set(style='whitegrid')

# 2. Load the Airbnb dataset
airbnbDataSet_filename = os.path.join(os.getcwd(), "data", "airbnbListingsData.
    ↪csv")
df = pd.read_csv(airbnbDataSet_filename)
print("Original shape:", df.shape)
display(df.head())

# 3. Basic cleaning

# Fill missing reviews_per_month with 0
if 'reviews_per_month' in df.columns:
    df['reviews_per_month'] = df['reviews_per_month'].fillna(0)

# Drop irrelevant columns if they exist
cols_to_drop = ['id', 'name', 'host_name', 'last_review']
existing_cols_to_drop = [col for col in cols_to_drop if col in df.columns]
df.drop(columns=existing_cols_to_drop, inplace=True)

# Remove outliers
df = df[(df['price'] <= 500) & (df['minimum_nights'] <= 365)]

# One-hot encode only the columns that exist
cols_to_encode = ['neighbourhood_group', 'room_type']
existing_cols = [col for col in cols_to_encode if col in df.columns]
df = pd.get_dummies(df, columns=existing_cols, drop_first=True)

print("Cleaned shape:", df.shape)
display(df.head())

# 4. Split data into features and label
X = df.drop('price', axis=1)
y = df['price']

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
    ↪random_state=42)

# Scale data (only for Linear Regression)
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# 5. Train and evaluate Linear Regression
lr = LinearRegression()
lr.fit(X_train_scaled, y_train)

```

```

y_pred_lr = lr.predict(X_test_scaled)

print("Linear Regression Performance:")
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred_lr)))
print("MAE:", mean_absolute_error(y_test, y_pred_lr))
print("R^2:", r2_score(y_test, y_pred_lr))

# 6. Train and evaluate Random Forest Regressor
rf = RandomForestRegressor(random_state=42)
rf.fit(X_train, y_train)
y_pred_rf = rf.predict(X_test)

print("\nRandom Forest Performance:")
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred_rf)))
print("MAE:", mean_absolute_error(y_test, y_pred_rf))
print("R^2:", r2_score(y_test, y_pred_rf))

# 7. Visualize feature importances
importances = rf.feature_importances_
features = X.columns
importance_df = pd.DataFrame({'Feature': features, 'Importance': importances})
importance_df.sort_values(by='Importance', ascending=False, inplace=True)

plt.figure(figsize=(10, 6))
sns.barplot(data=importance_df.head(10), x='Importance', y='Feature')
plt.title('Top 10 Feature Importances (Random Forest)')
plt.show()

# 8. Hyperparameter tuning for Random Forest
param_grid = {
    'n_estimators': [100, 200],
    'max_depth': [10, 20, None],
    'min_samples_split': [2, 5],
}

grid_search = GridSearchCV(RandomForestRegressor(random_state=42), param_grid,
                           cv=3, scoring='neg_root_mean_squared_error',
                           ↪n_jobs=-1)
grid_search.fit(X_train, y_train)
best_rf = grid_search.best_estimator_
y_pred_best = best_rf.predict(X_test)

print("\nTuned Random Forest Performance:")
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred_best)))
print("MAE:", mean_absolute_error(y_test, y_pred_best))
print("R^2:", r2_score(y_test, y_pred_best))

```

```
# 9. Final Summary
print("\nModel Comparison Summary")
print("-----")
print(f"Linear Regression R^2: {r2_score(y_test, y_pred_lr):.3f}")
print(f"Random Forest R^2: {r2_score(y_test, y_pred_rf):.3f}")
print(f"Tuned Random Forest R^2: {r2_score(y_test, y_pred_best):.3f}")
```

Original shape: (28022, 50)

	name \		description \	neighborhood_overview	host_name \	host_location \	host_about	host_response_rate \	host_acceptance_rate	host_is_superhost	host_listings_count	...
0	Skylit Midtown Castle		Beautiful, spacious skylit studio in the heart...	Centrally located in the heart of Manhattan ju...	Jennifer	New York, New York, United States	A New Yorker since 2000! My passion is creatin...	0.80	0.17	True	8.0	...
1	Whole flr w/private bdrm, bath & kitchen(pls r...		Enjoy 500 s.f. top floor in 1899 brownstone, w...	Just the right mix of urban center and local n...	LisaRoxanne	New York, New York, United States	Laid-back Native New Yorker (formerly bi-coast...	0.09	0.69	True	1.0	...
2	Spacious Brooklyn Duplex, Patio + Garden		We welcome you to stay in our lovely 2 br dupl...		Rebecca	Brooklyn, New York, United States	Rebecca is an artist/designer, and Henoch is i...	1.00	0.25	True	1.0	...
3	Large Furnished Room Near B'way		Please don't expect the luxury here just a bas...	Theater district, many restaurants around here.	Shunichi	New York, New York, United States	I used to work for a financial industry but no...	1.00				
4	Cozy Clean Guest Room - Family Apt		Our best guests are seeking a safe, clean, spa...	Our neighborhood is full of restaurants and ca...	MaryEllen	New York, New York, United States	Welcome to family life with my oldest two away...	NaN				

3	1.00	True	1.0 ...
4	NaN	True	1.0 ...

	review_scores_communication	review_scores_location	review_scores_value \
0	4.79	4.86	4.41
1	4.80	4.71	4.64
2	5.00	4.50	5.00
3	4.42	4.87	4.36
4	4.95	4.94	4.92

	instant_bookable	calculated_host_listings_count \
0	False	3
1	False	1
2	False	1
3	False	1
4	False	1

	calculated_host_listings_count_entire_homes \
0	3
1	1
2	1
3	0
4	0

	calculated_host_listings_count_private_rooms \
0	0
1	0
2	0
3	1
4	1

	calculated_host_listings_count_shared_rooms	reviews_per_month \
0	0	0.33
1	0	4.86
2	0	0.02
3	0	3.68
4	0	0.87

	n_host_verifications
0	9
1	6
2	3
3	4
4	7

[5 rows x 50 columns]

Cleaned shape: (27184, 50)

	description \
0	Beautiful, spacious skylit studio in the heart...
1	Enjoy 500 s.f. top floor in 1899 brownstone, w...
2	We welcome you to stay in our lovely 2 br dupl...
3	Please don't expect the luxury here just a bas...
4	Our best guests are seeking a safe, clean, spa...

	neighborhood_overview \
0	Centrally located in the heart of Manhattan ju...
1	Just the right mix of urban center and local n...
2	NaN
3	Theater district, many restaurants around here.
4	Our neighborhood is full of restaurants and ca...

	host_location \
0	New York, New York, United States
1	New York, New York, United States
2	Brooklyn, New York, United States
3	New York, New York, United States
4	New York, New York, United States

	host_about	host_response_rate \
0	A New Yorker since 2000! My passion is creatin...	0.80
1	Laid-back Native New Yorker (formerly bi-coast...	0.09
2	Rebecca is an artist/designer, and Henoch is i...	1.00
3	I used to work for a financial industry but no...	1.00
4	Welcome to family life with my oldest two away...	NaN

	host_acceptance_rate	host_is_superhost	host_listings_count \
0	0.17	True	8.0
1	0.69	True	1.0
2	0.25	True	1.0
3	1.00	True	1.0
4	NaN	True	1.0

	host_total_listings_count	host_has_profile_pic	...	instant_bookable \
0	8.0	True	...	False
1	1.0	True	...	False
2	1.0	True	...	False
3	1.0	True	...	False
4	1.0	True	...	False

	calculated_host_listings_count	calculated_host_listings_count_entire_homes \
0	3	3
1	1	1
2	1	1

3	1	0
4	1	0

	calculated_host_listings_count_private_rooms \
0	0
1	0
2	0
3	1
4	1

	calculated_host_listings_count_shared_rooms	reviews_per_month \
0	0	0.33
1	0	4.86
2	0	0.02
3	0	3.68
4	0	0.87

	n_host_verifications	room_type_Hotel room	room_type_Private room \
0	9	0	0
1	6	0	0
2	3	0	0
3	4	0	1
4	7	0	1

	room_type_Shared room
0	0
1	0
2	0
3	0
4	0

[5 rows x 50 columns]

↳

↳ last) ValueError Traceback (most recent call↳

```

<ipython-input-7-167920ad036d> in <module>()
    52 # Scale data (only for Linear Regression)
    53 scaler = StandardScaler()
----> 54 X_train_scaled = scaler.fit_transform(X_train)
    55 X_test_scaled = scaler.transform(X_test)
    56

```

```

/usr/local/lib/python3.6/dist-packages/sklearn/base.py in
fit_transform(self, X, y, **fit_params)
    569         if y is None:
    570             # fit method of arity 1 (unsupervised transformation)
--> 571             return self.fit(X, **fit_params).transform(X)
    572         else:
    573             # fit method of arity 2 (supervised transformation)

/usr/local/lib/python3.6/dist-packages/sklearn/preprocessing/_data.py in
fit(self, X, y)
    667         # Reset internal state before fitting
    668         self._reset()
--> 669         return self.partial_fit(X, y)
    670
    671     def partial_fit(self, X, y=None):

/usr/local/lib/python3.6/dist-packages/sklearn/preprocessing/_data.py in
partial_fit(self, X, y)
    698         X = check_array(X, accept_sparse=('csr', 'csc'),
    699                           estimator=self, dtype=FLOAT_DTYPES,
--> 700                           force_all_finite='allow-nan')
    701
    702         # Even in the case of `with_mean=False`, we update the mean
anyway

/usr/local/lib/python3.6/dist-packages/sklearn/utils/validation.py in
check_array(array, accept_sparse, accept_large_sparse, dtype, order, copy,
force_all_finite, ensure_2d, allow_nd, ensure_min_samples,
ensure_min_features, warn_on_dtype, estimator)
    529         array = array.astype(dtype, casting="unsafe",
copy=False)
    530         else:
--> 531         array = np.asarray(array, order=order,
dtype=dtype)
    532         except ComplexWarning:
    533             raise ValueError("Complex data not supported\n"

~/local/lib/python3.6/site-packages/numpy/core/_asarray.py in
asarray(a, dtype, order)
    81
    82     """
---> 83     return array(a, dtype, copy=False, order=order)

```


84
85

```
~/.local/lib/python3.6/site-packages/pandas/core/generic.py in
↳ __array__(self, dtype)
    1779
    1780     def __array__(self, dtype=None) -> np.ndarray:
-> 1781         return np.asarray(self._values, dtype=dtype)
    1782
    1783     def __array_wrap__(self, result, context=None):

~/.local/lib/python3.6/site-packages/numpy/core/_asarray.py in
↳ asarray(a, dtype, order)
    81
    82     """
---> 83     return array(a, dtype, copy=False, order=order)
    84
    85
```

ValueError: could not convert string to float: "Fantastic Hudson Yards 3
↳ Bed 1 Bath located just a short walk to Times Square 42nd Street. Amazing
↳ access to all that makes NYC great. Just two flights in a traditional NYC walk
↳ up building. Beautiful character throughout with original brick wall and
↳ hardwood accents. Two queen beds and one full each with space to store
↳ belongings. High Speed WiFi and Flat Screen TV with Netflix provided. Dining
↳ table seats 6 and kitchen has all you'll need to prepare meal at home."